

Elastic Load Balancer and Auto Scaling

Assignment – 1 ELB, ASG and 53

Problem Statement

You work for XYZ Corporation that uses on premise solutions and a limited number of systems. With the increase in request in their application, the load also increases. So, to handle the load the corporation has to buy more system almost on a regular basis. Realizing the need to cut down the expenses on systems, they decide to move their infrastructure to AWS.

Task to be performed:

1. Manage the scaling requirements of the company by:
 - a. Deploying multiple compute resources on the cloud as soon as the load increases and the CPU utilization exceeds 80%
 - b. Removing the resource when the CPU utilization goes under 60%
2. Create a load balancer to distribute the load between compute resources.
3. Route the traffic to the company's domain.

Aim

- Create an instance → Install webserver → Custom webserver
- Create an AMI out of this instance
- Launch configuration/template
- Create Auto scaling group
- Create Dynamic scaling policy
 - a. CPU utilization more than 80%
 - b. CPU utilization less than 60%
- Create a load balancer
- Finally route the traffic to port 53

Steps

- Create a EC2 instance and connect to instance
- Update the instance and also install webserver and customize the web server
- Select the instance and Go to action
- Select image and templates and select create image

- Create an AMI image

The screenshot shows the AWS CloudShell interface for launching an instance. The main area displays the 'Name and tags' section with a 'Name' field containing 'Task'. Below it is the 'Application and OS Images (Amazon Machine Image)' section, which lists the selected AMI as 'Canonical, Ubuntu, 22.04 LTS, ami-0c7217cdde317cfec'. Configuration details include a 'Virtual server type (instance type)' of 't2.micro', a 'Firewall (security group)' of 'New security group', and a 'Storage (volumes)' of '1 volume(s) ~ 8 GiB'. On the right, a summary panel shows 'Number of instances' set to 1. A prominent orange 'Launch instance' button is located at the bottom right of the configuration area.

The screenshot shows the AWS CloudShell terminal window. The output of the command 'sudo apt update' is displayed, indicating that the list of available updates is more than a week old and prompting to run 'sudo apt update' to check for new updates. It also shows the standard Ubuntu ESM (Enterprise Software Maintenance) notice about software being free and distribution terms. The terminal prompt shows the instance ID 'i-024d67e6ae24af9e' and its public IP '3.87.205.150'.

```

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

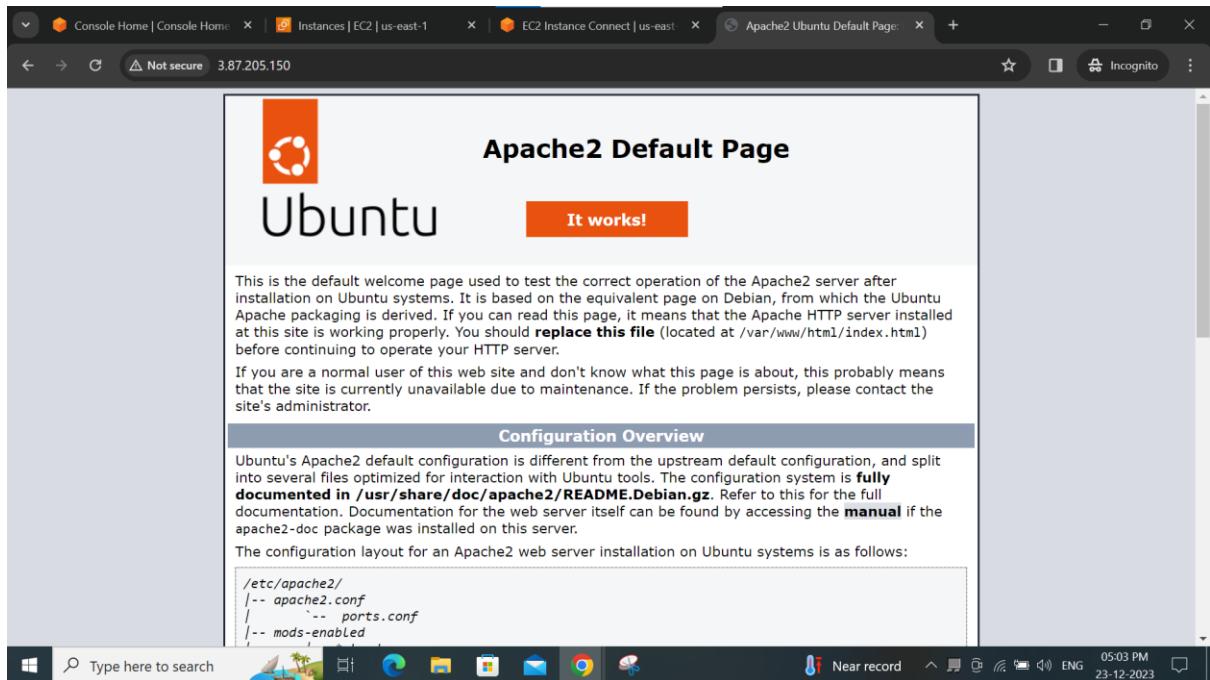
The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

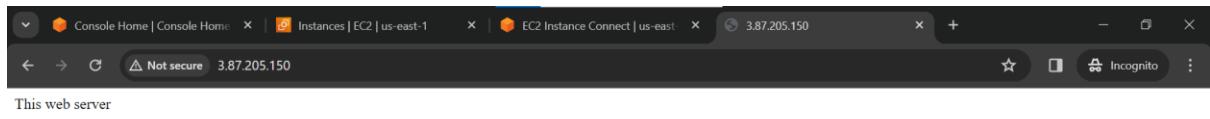
ubuntu@ip-172-31-96-100:~$
```



```
ubuntu@ip-172-31-86-100:/var/www/html$ ls
index.html
ubuntu@ip-172-31-86-100:/var/www/html$ sudo rm index.html
ubuntu@ip-172-31-86-100:/var/www/html$ ls
ubuntu@ip-172-31-86-100:/var/www/html$ sudo nano index.html[]
```

i-024d67e6ae24af9e (Task)
Public IPs: 3.87.205.150 Private IPs: 172.31.86.100

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences



The screenshot shows two pages from the AWS EC2 service.

Create image (Info)

An image (also referred to as an AMI) defines the programs and settings that are applied when you launch an EC2 instance. You can create an image from the configuration of an existing instance.

Instance ID: i-024d67e6ae24af9e (Task)

Image name: My task

Image description - optional: Image description

No reboot: Enable

Instance volumes: (Listed as empty)

Actions: CloudShell, Feedback

Amazon Machine Images (AMIs) (1) Info

Owned by me: my task AMI

Name	AMI name	AMI ID	Source
my task AMI	My task AMI	ami-0e19fc16bd7bd27c8	237981402912/My task AMI

Select an AMI

© 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

- Create auto scaling group
- Create a launch template
 - a. Select the AMI which was created
 - b. And add necessary configuration and create a template
- Select the launch template which was created
- Map to at least one availability zone where EC2 instance is created
- Provide the Group size Desired (capacity - 4)
- Provide minimum and maximum capacity (minimum capacity - 1 and maximum capacity - 8)
- Once the auto scaling group is created 4 EC2 instance will be created by the auto scaling group because the desire capacity is 4
- If we put the public ID address of 4 new created instance in browser the customized web page same of EC2 instance which was created will be reflected

The screenshot shows the AWS CloudFormation Create Stack Wizard. The current step is "Step 1 Choose stack name". The "Name" field contains the value "task". A note below the field says: "Must be unique to this account in the current Region and no more than 255 characters." On the left, there's a sidebar with navigation links: "EC2 > Auto Scaling groups > Create Auto Scaling group", "Choose stack name", "Configure advanced options", "Configure group size and scaling", "Add notifications", and "Add tags". At the bottom, there are links for "CloudShell", "Feedback", and "Cookie preferences".

Creating a launch template allows you to create a saved instance configuration that can be reused, shared and launched at a later time. Templates can have multiple versions.

Launch template name and description

Launch template name - required

Must be unique to this account. Max 128 chars. No spaces or special characters like '&', '*', '@'.

Template version description

Max 255 chars

Auto Scaling guidance [Info](#)

Select this if you intend to use this template with EC2 Auto Scaling

Provide guidance to help me set up a template that I can use with EC2 Auto Scaling

Template tags

Source template

Summary

Software Image (AMI)

Virtual server type (instance type)

Firewall (security group)

Storage (volumes)

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in

Create launch template

Services [Alt+S] N. Virginia Naveen

Don't include in launch template Owned by me Shared with me

Browse more AMIs Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

My task AMI ami-0e19fc16bd7bd27c8 2023-12-23T11:40:36.000Z Virtualization: hvm ENA enabled: true Root device type: ebs

Description
My task AMI

Architecture x86_64 **AMI ID** ami-0e19fc16bd7bd27c8

Instance type [Info](#) | [Get advice](#) Advanced

Create launch template

Configure advanced options

Step 4 - optional Configure group size and scaling

Step 5 - optional Add notifications

Step 6 - optional Add tags

Step 7 Review

Auto Scaling group name Enter a name to identify the group.

Must be unique to this account in the current Region and no more than 255 characters.

Launch template [Info](#)

For accounts created after May 31, 2023, the EC2 console only supports creating Auto Scaling groups with launch templates. Creating Auto Scaling groups with launch configurations is not recommended but still available via the CLI and API until December 31, 2023.

Launch template Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

Create a launch template

Version Default (1)

Step 6 - optional

[Add tags](#)

Step 7

[Review](#)

Network [Info](#)

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-0d30cb522f7687a00
172.31.0.0/16 Default

[Create a VPC](#)

Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets

us-east-1a | subnet-03639ccc1cfb5291a X
172.31.32.0/20 Default

us-east-1b | subnet-09de3c6cd1daac31f X
172.31.0.0/20 Default

us-east-1c | subnet-0525e8c3686e7911a X
172.31.80.0/20 Default

[Create a subnet](#)

© 2023, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

Step 5 - optional

[Add notifications](#)

Step 6 - optional

[Add tags](#)

Step 7

[Review](#)

Units (number of instances)

Desired capacity

Specify your group size.

4

Scaling [Info](#)

You can resize your Auto Scaling group manually or automatically to meet changes in demand.

Scaling limits

Set limits on how much your desired capacity can be increased or decreased.

Min desired capacity Max desired capacity

1 8

Equal or less than desired capacity Equal or greater than desired capacity

Automatic scaling - optional

Choose whether to use a target tracking policy [Info](#)

You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.

No scaling policies Your Auto Scaling group will remain at its initial size and

Target tracking scaling policy Choose a CloudWatch metric and target value and let the

© 2023, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

Step 5: Add notifications

[Edit](#)

Notifications

No notifications

Step 6: Add tags

[Edit](#)

Tags (1)

Key	Value	Tag new instances
name	my ec2 auto scaling	Yes

[Cancel](#) [Previous](#) [Create Auto Scaling group](#)

© 2023, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

AWS Services Search [Alt+S] N. Virginia Naveen

Activity history (4)

Filter activity history

Status	Description	Cause	Start time
Successful	Launching a new EC2 instance: i-0c9d6e750322e40ad	At 2023-12-23T12:04:16Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 4. At 2023-12-23T12:04:28Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 4.	2023 December 23, 05:34:28 PM +05:30
Successful	Launching a new EC2 instance: i-0e1f1d91a69029814	At 2023-12-23T12:04:16Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 4. At 2023-12-23T12:04:28Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 4.	2023 December 23, 05:34:28 PM +05:30
Successful	Launching a new EC2 instance: i-0c9d901bcf9d7eefa	At 2023-12-23T12:04:16Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 4. At 2023-12-23T12:04:28Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 4.	2023 December 23, 05:34:28 PM +05:30
Successful	Launching a new EC2	At 2023-12-23T12:04:16Z a user request created an AutoScalingGroup changing the desired capacity	2023 December 23, 05:34:28 PM +05:30

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

AWS Services Search [Alt+S] N. Virginia Naveen

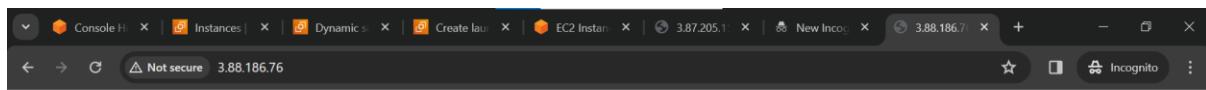
Instances (5) Info

Find Instance by attribute or tag (case-sensitive)

Name	Instance ID	Instance state	Instance type	Status check	Alarm status
Task	i-024d67e6ae24af9e	Running	t2.micro	2/2 checks passed	No alarms
	i-0c9d6e750322e40ad	Running	t2.micro	Initializing	No alarms
	i-0e1f1d91a69029814	Running	t2.micro	2/2 checks passed	No alarms
	i-0c9d901bcf9d7eefa	Running	t2.micro	2/2 checks passed	No alarms

Select an instance

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences



Create Dynamic scaling policy CPU utilization more than 80%

- Go inside the auto scaling group and click automatic scaling
- Now click create dynamic scaling policy and select policy type as step scaling
- Create a cloud watch alarm
 - a) Select the metric as CPU utilization
 - b) Select as Greater and provide the threshold value as 80%
 - c) Provide an alarm name
- Select the cloud watch alarm which was created
- Select ADD in take the action as 1 (Increase the instance one by one whenever CUP utilization is more than 80%)
- And Click create

Create Dynamic scaling policy CPU utilization more than 80%

- Go inside the auto scaling group and click automatic scaling
- Now click create dynamic scaling policy and select policy type as step scaling
- Create a cloud watch alarm
 - d) Select the metric as CPU utilization
 - e) Select as Lower and provide the threshold value as 60%
 - f) Provide an alarm name
- Select the cloud watch alarm which was created
- Select REMOVE in take the action as and provide 1 (Decrease the instance one by one whenever CUP utilization is less than 60%)
- And Click create
- Once this policy is created the EC2 instance will be terminate one by one and remain only one instance which is created by auto scaling group and EC2 instance which is created (since the minimum capacity – 1) this happens because currently CUP utilization is under 60%

AWS Services Search [Alt+S] N. Virginia Naveen

EC2 > Auto Scaling groups > task

task

Details Activity Automatic scaling Instance management Monitoring Instance refresh

Scaling policies resize your Auto Scaling group to meet changes in demand. With reactive dynamic scaling policies, you can track specific CloudWatch metrics and take action when the CloudWatch alarm threshold is met. Use predictive scaling policies along with dynamic scaling policies in the following situations: when your application demand changes quickly, but with a recurring pattern, or when your EC2 instances require more time to initialize.

Dynamic scaling policies (0) info Actions Create dynamic scaling policy < 1 >

No dynamic scaling policies have been created

Dynamic scaling policies use real-time data to scale your group based on configurable metrics.

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

AWS Services Search [Alt+S] N. Virginia Naveen

CloudWatch Metrics Step 1: Specify Step 2: Configure Step 3: Add name Step 4: Preview

Select metric

Untitled graph 1h 3h 12h 1d 3d 1w Custom UTC timezone Line C

Browse Multi source query - new Graphed metrics (1) Options Source Add math Add query

task	NetworkPacketsIn	No alarms
<input checked="" type="checkbox"/> task	CPUUtilization	No alarms
task	NetworkIn	No alarms
task	CPUCreditBalance	No alarms
task	CPUCreditUsage	No alarms

Select metric

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

AWS Services Search [Alt+S] N. Virginia Naveen

Step 2: Configure actions Step 3: Add name and description Step 4: Preview and create

Metric

Graph This alarm will trigger when the blue line goes above the red line for 1 datapoints within 1 minute.

Percent

80 80

40.1

0.131 09:30 10:30 11:30

CPUUtilization

Edit Namespace AWS/EC2 Metric name CPUUtilization AutoScalingGroupName task Statistic Average Period 1 minute

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

AWS Services Search [Alt+S] N. Virginia Naveen

Threshold type

Static Use a value as a threshold

Anomaly detection Use a band as a threshold

Whenever CPUUtilization is...

Greater > threshold

Greater/Equal >= threshold

Lower/Equal <= threshold

Lower < threshold

than... Define the threshold value.

80 Must be a number

▶ Additional configuration

Cancel Next

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

CloudWatch > Alarms > Create alarm Step 1 Specify metric and conditions

Step 2 Configure actions

Step 3 Add name and description

Step 4 Preview and create

Add name and description

Name and description

Alarm name

Alarm description - optional [View formatting guidelines](#)

Edit Preview

This is an H1
double asterisks will produce strong character
This is [an example](https://example.com/) inline link.

Up to 1024 characters (0/1024)

Markdown formatting is only applied when viewing your alarm in the console. The description will remain in

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

EC2 > Auto Scaling groups > task

Create dynamic scaling policy

Policy type

Scaling policy name

CloudWatch alarm Choose an alarm that can scale capacity whenever:

alarm-cpu-utilization-80

Create a CloudWatch alarm

breaches the alarm threshold: CPUUtilization > 80 for 1 consecutive periods of 60 seconds for the metric dimensions:

AutoScalingGroupName = task

Take the action

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Choose an alarm that can scale capacity whenever:

alarm-cpu-utilization-80

Create a CloudWatch alarm

breaches the alarm threshold: CPUUtilization > 80 for 1 consecutive periods of 60 seconds for the metric dimensions:

AutoScalingGroupName = task

Take the action

Add

1 capacity units when 80 <= CPUUtilization < +infinity

Add step

Instance warmup 60 seconds

Cancel **Create**

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

aws Services Search [Alt+S]

Threshold type

Static Use a value as a threshold

Anomaly detection Use a band as a threshold

Whenever CPUUtilization is...

Define the alarm condition.

Greater > threshold

Greater/Equal >= threshold

Lower/Equal <= threshold

Lower < threshold

than...

Define the threshold value.

60

Must be a number

Additional configuration

Cancel **Next**

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

aws Services Search [Alt+S]

Step 1: Specify metric and conditions

Step 2: Configure actions

Step 3: Add name and description

Step 4: Preview and create

Add name and description

Name and description

Alarm name

alarm-cpu-utilization-60

Alarm description - optional [View formatting guidelines](#)

Edit Preview

This is an H1
double asterisks will produce strong character
This is [an example](https://example.com/) inline link.

Up to 1024 characters (0/1024)

Markdown formatting is only applied when viewing your alarm in the console. The description will remain in plain text in the alarm notifications.

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

AWS Services Search [Alt+S] N. Virginia Naveen

EC2 > Auto Scaling groups > task

Create dynamic scaling policy

Policy type: Step scaling

Scaling policy name: cpu-utilization-60-policy

CloudWatch alarm: Choose an alarm that can scale capacity whenever: alarm-cpu-utilization-60

breaches the alarm threshold: CPUUtilization < 60 for 1 consecutive periods of 300 seconds for the metric dimensions:

AutoScalingGroupName = task

Take the action: Remove

1 capacity units when 60 >= CPUUtilization > -infinity

Add step

Cancel Create

© 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

AWS Services Search [Alt+S] N. Virginia Naveen

cpu-utilization-60-policy

CloudWatch alarm: Choose an alarm that can scale capacity whenever: alarm-cpu-utilization-60

breaches the alarm threshold: CPUUtilization < 60 for 1 consecutive periods of 300 seconds for the metric dimensions:

AutoScalingGroupName = task

Take the action: Remove

1 capacity units when 60 >= CPUUtilization > -infinity

Add step

Cancel Create

© 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences 06:00 PM

AWS Services Search [Alt+S] N. Virginia Naveen

Dynamic scaling policy created or edited successfully.

cpu-utilization-60-policy Step scaling Enabled alarm-cpu-utilization-60 breaches the alarm threshold: CPUUtilization < 60 for 1 consecutive periods of 300 seconds for the metric dimensions: AutoScalingGroupName = task Remove 1 capacity units when 60 >= CPUUtilization > -infinity

cpu-utilization-80-policy Step scaling Enabled alarm-cpu-utilization-80 breaches the alarm threshold: CPUUtilization > 80 for 1 consecutive periods of 60 seconds for the metric dimensions: AutoScalingGroupName = task Add 1 capacity units when 80 <= CPUUtilization < +infinity 60 seconds to warm up after each step

Predictive scaling policies (0) Info Actions Create predictive scaling policy

© 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

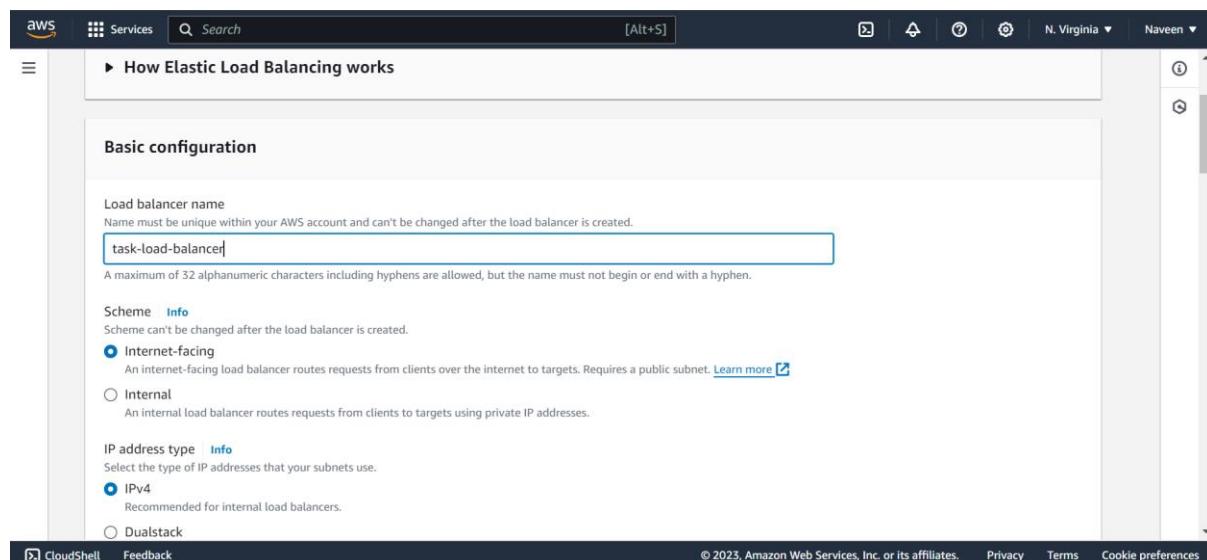
[Alt+S]			
	Status	Description	Cause
	Successful	Terminating EC2 instance: i-0c9d901bcf9d7eefaa	At 2023-12-23T12:32:37Z a monitor alarm alarm-cpu-utilization-60 in state ALARM triggered cpu-utilization-60-policy changing the desired capacity from 4 to 3. At 2023-12-23T12:33:37Z instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 4 to 3. At 2023-12-23T12:32:44Z instance i-0c9d901bcf9d7eefaa selected for termination.
	Successful	Launching a new EC2 instance: i-0c9d6e750322e40ad	At 2023-12-23T12:04:16Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 4. At 2023-12-23T12:04:28Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 4.
	Successful	Launching a new EC2 instance: i-0e1f1d91a69029814	At 2023-12-23T12:04:16Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 4. At 2023-12-23T12:04:28Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 4.
	Successful	Launching a new EC2 instance: i-0c9d901bcf9d7eefaa	At 2023-12-23T12:04:16Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 4. At 2023-12-23T12:04:28Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 4.

[Alt+S]			
	Name	Instance ID	Instance state
	Task	i-024d67e6ae24af9e	Running
		i-0c9d6e750322e40ad	Shutting-down
		i-0e1f1d91a69029814	Running
		i-0c9d901bcf9d7eefaa	Terminated
		i-0e5f2e85d4981464e	Terminated

[Alt+S]			
	Status	Description	Cause
	Successful	Terminating EC2 instance: i-0c9d6e750322e40ad	At 2023-12-23T12:34:37Z a monitor alarm alarm-cpu-utilization-60 in state ALARM triggered cpu-utilization-60-policy changing the desired capacity from 2 to 1. At 2023-12-23T12:34:37Z instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 2 to 1. At 2023-12-23T12:34:44Z instance i-0c9d6e750322e40ad selected for termination.
	Successful	Terminating EC2 instance: i-0e5f2e85d4981464e	At 2023-12-23T12:33:37Z a monitor alarm alarm-cpu-utilization-60 in state ALARM triggered cpu-utilization-60-policy changing the desired capacity from 3 to 2. At 2023-12-23T12:33:37Z instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 3 to 2. At 2023-12-23T12:33:38Z instance i-0e5f2e85d4981464e selected for termination.
	Successful	Terminating EC2 instance: i-0c9d901bcf9d7eefaa	At 2023-12-23T12:32:37Z a monitor alarm alarm-cpu-utilization-60 in state ALARM triggered cpu-utilization-60-policy changing the desired capacity from 4 to 3. At 2023-12-23T12:32:37Z instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 4 to 3. At 2023-12-23T12:32:44Z instance i-0c9d901bcf9d7eefaa selected for termination.
	Successful	Launching a new EC2 instance: i-0c9d6e750322e40ad	At 2023-12-23T12:04:16Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 4. At 2023-12-23T12:04:28Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 4.

Create a load balancer to distribute the load between compute resources

- Create a load application load balancer and provide a name to the load balancer
- Create a target group
 - a) Choose a target type as instance
 - b) And provide a name to target
 - c) Next protocol is http → port is 80 and IPV4
 - d) Register the target by selecting the EC2 instance
 - e) And Click a target group
- Select the target group which was created in listener and routing
- And click a create load balancer
- Now if we put the DNS name in the Brower the customized web page same of EC2 instance which was created will be reflected



AWS Services Search [Alt+S] N. Virginia Naveen

EC2 > Target groups > Create target group

Step 1 Specify group details

Step 2 Register targets

Specify group details

Your load balancer routes requests to the targets in a target group and performs health checks on the targets.

Basic configuration

Settings in this section can't be changed after the target group is created.

Choose a target type

Instances

- Supports load balancing to instances within a specific VPC.
- Facilitates the use of [Amazon EC2 Auto Scaling](#) to manage and scale your EC2 capacity.

IP addresses

- Supports load balancing to VPC and on-premises resources.
- Facilitates routing to multiple IP addresses and network interfaces on the same instance.
- Offers flexibility with microservice based architectures, simplifying inter-application communication.
- Supports IPv6 targets, enabling end-to-end IPv6 communication, and IPv4-to-IPv6 NAT.

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

AWS Services Search [Alt+S] N. Virginia Naveen

Accessible to Application Load Balancers only.

Application Load Balancer

- Offers the flexibility for a Network Load Balancer to accept and route TCP requests within a specific VPC.
- Facilitates using static IP addresses and PrivateLink with an Application Load Balancer.

Target group name

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Protocol : Port

Choose a protocol for your target group that corresponds to the Load Balancer type that will route traffic to it. Some protocols now include anomaly detection for the targets and you can set mitigation options once your target group is created. This choice cannot be changed after creation

HTTP 80 1-65535

IP address type

Only targets with the indicated IP address type can be registered to this target group.

IPv4

Each instance has a default network interface (eth0) that is assigned the primary private IPv4 address. The instance's primary private IPv4 address is the one that will be applied to the target.

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

AWS Services Search [Alt+S] N. Virginia Naveen

Mappings Info

Select at least two Availability Zones and one subnet per zone. The load balancer routes traffic to targets in these Availability Zones only. Availability Zones that are not supported by the load balancer or the VPC are not available for selection.

us-east-1a (use1-az6)

Subnet

IPv4 address

Assigned by AWS

us-east-1b (use1-az1)

Subnet

IPv4 address

Assigned by AWS

us-east-1c (use1-az2)

Subnet

IPv4 address

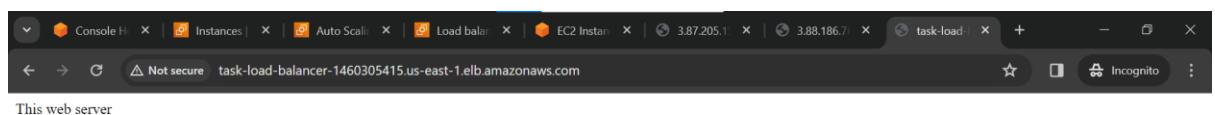
CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

The screenshot shows the AWS CloudFront Listener configuration page. A listener is defined for port 80, forwarding to a target group named 'task'. The target type is set to Instance, IPv4. There are sections for 'Listener tags - optional' and 'Add-on services - optional'.

Protocol: HTTP
Port: 80
Default action: task
Create target group

Listener tags - optional
Consider adding tags to your listener. Tags enable you to categorize your AWS resources so you can more easily manage them.
Add listener tag
You can add up to 50 more tags.

Add-on services - optional
Additional AWS services can be integrated with this load balancer at launch. You can also add these and other services after your load balancer is created by reviewing the "Integrated Services" tab for the selected load balancer.



Route the traffic to the company's domain

- Click the register domain name
- Once domain register click the hosted zone
- And provide the domain name
- And click create hosted zone
- Go inside the hosted zone and create new record
- Provide record name and select alias
- And select alias to application and classic load balancer
- Select the load balancer which is created
- And click create record
- Now the route traffic to company domain (**Register domain name was not done since the registering the domain name is costly**)

The screenshot shows the AWS Route 53 'Create record' interface. At the top, there are two main options: 'Quick create (recommended for expert users)' and 'Wizard (recommended for new users)'. The 'Quick create' section is selected, showing fields for 'Record name' (set to 'my') and 'Record type' (set to 'A – Routes traffic to an IPv4 address and some AWS resources'). Below these are checkboxes for 'Alias' (selected) and 'Route traffic to' (set to 'Info'). The 'Wizard' section is described as 'Choose this method if you need more explanations as you create your record.' At the bottom of the page, there are links for 'CloudShell', 'Feedback', and copyright information ('© 2023, Amazon Web Services, Inc. or its affiliates.'), along with standard navigation buttons like 'Cancel' and 'Create records'.

Quick create (recommended for expert users)

Wizard (recommended for new users)

Choose this method if you are confident in the process of creating records and know which options you need.

Choose this method if you need more explanations as you create your record.

Create record [Info](#)

Quick create record [Switch to wizard](#)

Record 1

Record name [Info](#) my .task.com

Record type [Info](#) A – Routes traffic to an IPv4 address and some AWS resources

Keep blank to create a record for the root domain.

Alias

Route traffic to [Info](#)

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Record 1

Record name [Info](#) my .task.com

Record type [Info](#) A – Routes traffic to an IPv4 address and some AWS resources

Keep blank to create a record for the root domain.

Alias

Route traffic to [Info](#)

Alias to Application and Classic Load Balancer

US East (N. Virginia)

Choose load balancer

Routing policy [Info](#) Simple routing

Evaluate target health Yes

Add another record

Cancel [Create records](#)

https://us-east-1.console.aws.amazon.com/console/home?region=us-east-1

© 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

AWS Services Search [Alt+S] Global Naveen

Record 1

Record name [Info](#) my .task.com Record type [Info](#) A – Routes traffic to an IPv4 address and some AWS resources

Keep blank to create a record for the root domain.

Alias

Route traffic to [Info](#) Alias to Application and Classic Load Balancer

US East (N. Virginia)

dualstack.task-load-balancer-1460305415.us-east-1.elb.amazonaws.com

Alias hosted zone ID: Z355XD0TRQ7XK

Routing policy [Info](#) Simple routing Evaluate target health Yes

Add another record

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

aws Services Search [Alt+S] Global Naveen

Route 53

Records (3) DNSSEC signing Hosted zone tags (0)

Records (3) [Info](#)

Automatic mode is the current search behavior optimized for best filter results. To change modes go to settings.

Delete record Import zone file Create record

Filter records by property or value Type Routing policy Alias

	Record ...	Type	Routing policy	Alias	Value/Route traffic to
<input type="checkbox"/>	task.com	NS	Simple	-	No ns-1774.awsdns-29.co.uk. ns-501.awsdns-62.com. ns-643.awsdns-16.net. ns-1185.awsdns-20.org.
<input type="checkbox"/>	task.com	SOA	Simple	-	No ns-1774.awsdns-29.co.uk. a...
<input type="checkbox"/>	my.task.com	A	Simple	-	Yes dualstack.task-load-balancer...

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences